

CUSTOMER DISCOVERY

LOG 4 — PROBLEM VALIDATION

Interview Guide — Cyber Defense Researchers

Venture:	ChelysSecurity	Format:	25–30 min, 1 interviewer + 1 note-taker if possible
Segment:	Cyber defense researchers — academic, vendor, government, independent	Record:	Audio with permission OR written notes immediately after
Goal:	Understand how researchers discover and validate attack techniques, where their workflow breaks down, and what they would need to trust an AI system making autonomous offensive decisions.		

BEFORE YOU START

Do not show the product or mention pricing until Section 5. Focus just on this interview to be honest and to pitch if they ask what you are building. Say: “I’ll show you at the end — I want to hear your experience first and don’t bias your answers. Every leading question you ask compromises your evidence.”

Section 1 — Background (3 min)

Warm-up only. Orient yourself to their role and don't spend more than 3 minutes here. The goal is to know whether they're offensive, defensive, or mixed — so you can tailor follow-up questions in Section 2.

Q1 What does your research focus on right now — are you more on the offensive side, detection engineering, threat intel, or something else?

- If they say 'a mix,' ask which one takes up most of their actual hands-on time.

Q2 What does a typical week look like for you in terms of tooling and hands-on work?

- Listen for tool names. Metasploit, Cobalt Strike, Caldera, Splunk, Elastic — these tell you their technical depth and existing workflow.

Q3 Have you ever been involved in red team exercises or purple team engagements, or is your work primarily defensive?

- Purple team involvement means they've seen both sides — their opinions on automation will be more nuanced and more useful.

Section 2 — Current Workflow and Pain Points (10 min)

This is the most important section. Do not rush it. Listen for frustration, manual workarounds, and repeated problems — those are the signals that map directly to product value.

Q4 When you're trying to understand how a specific attack technique works in a real environment — how do you go about that today?

- If they describe a manual process (standing up a lab, running individual tool commands, writing custom scripts), ask: 'How long does that typically take you?'

Q5 What tools do you rely on most heavily for that kind of testing or emulation?

- Note every tool name. Ask: 'Is there anything about [tool] that frustrates you or that you wish worked differently?'

Q6 Where does that process break down or slow you down the most?

- This is your primary pain discovery question. If they pause, wait. Do not fill the silence. The first answer is often surface-level; the real pain comes after a few seconds of thinking.

Q7 How do you validate that a defensive control actually catches what it's supposed to catch? Walk me through a recent example if you can.

- A concrete example is worth 10 general statements. If they give one, ask: 'What was the hardest part of that specific test?' If they can't give one, ask: 'What's the closest example you can think of?'

Q8 How often do you feel like you're testing against the same known techniques over and over versus finding something genuinely new in your environment?

- High repetition = opportunity for automation. Desire for novelty = opportunity for AI-driven path discovery. Both map to ChelysSecurity features — listen for which one resonates more.

Q9 When you hand off findings or recommendations to a blue team or engineering team — what usually gets lost or ignored, and why do you think that happens?

- This surfaces the remediation gap — whether findings are actionable, whether the output format matters, and whether there's a communication problem that goes beyond just finding the vulnerability.

Section 3 — Attitudes Toward Automation and AI (8 min)

Do not mention your product yet. These questions surface their existing mental model of AI in security so you know what objections to expect and what language they use. If they react negatively, that is valuable data — do not defend AI, just ask why.

Q1 0 Have you used any automated tools for attack simulation or emulation — things like Atomic Red Team, Caldera, or breach-and-attack simulation platforms? What was your experience?

- If yes: 'What did it get right versus what a skilled human attacker would do?' If no: 'What stopped you from trying them?'

Q1 1 What do those tools get right? What do they get wrong compared to what a skilled human attacker would actually do?

- The gap between automated tools and skilled humans is exactly the product opportunity. Quote whatever they say here verbatim — it's your strongest evidence.

Q1 2 If you imagine a fully autonomous system that could pick its own attack path through an environment based on what it discovers — what's your gut reaction to that? What excites you, what makes you skeptical?

- After they answer, ask: 'Say more about that.' Researchers are precise thinkers — their second answer is usually more revealing than their first. Do not jump in with reassurances about safety controls.

Q1 3 What would that system need to do correctly for you to trust its findings enough to act on them?

- This is your trust threshold question. Listen for: reproducibility, explainability, MITRE mapping accuracy, false positive rate, evidence quality, scope adherence. Each one is a feature validation.

Q1 4 Where do you think AI makes security research better right now, and where do you think it's being oversold?

- Let them be critical. If they say 'AI pentesting is a gimmick,' ask: 'What would change your mind?' That answer is more useful than agreement.

Section 4 — Technical Credibility and Trust (5 min)

These questions are specifically about what it would take for a researcher to use or vouch for a tool that makes autonomous offensive decisions. The answers tell you what your Safety Oracle, audit trail, and finding output need to look like for this segment to take it seriously.

Q1 5 If an AI system told you it found a critical vulnerability using a specific technique — what would you need to see in the output before you believed it?

- Listen for: raw evidence (packets, screenshots, request/response), reproducible steps, CVE reference, MITRE technique ID, confidence score. These map directly to our finding data model.

Q1 6 How important is MITRE ATT&CK coverage to you personally when evaluating a tool? Is it a meaningful signal or mostly checkbox compliance theater?

- Researchers who say 'checkbox theater' are telling you they want technique-level accuracy, not just tactic labels. That's a higher bar — and a stronger differentiator if we meet it.

Q1 7 What's the worst outcome from a false positive in your research context — a missed finding, a wasted investigation, something else?

- This tells you how much false positive tolerance they have. High tolerance = they'll accept AI output as a starting point. Zero tolerance = they need verified, reproducible evidence for every finding.

Q1 8 If a tool made an autonomous decision during a test that you hadn't explicitly authorized — how would you feel about that, and what safeguards would you want in place?

- This is the Safety Oracle validation question. If they say 'that would be a dealbreaker without hard controls' — write that quote down. If they say 'I'd want a full audit log of every decision' — that's also a direct feature signal.

Section 5 — Closing (3 min)

End every interview with these three questions. Q19 surfaces gaps your product can own. Q20 is your referral engine — one researcher who connects you to two colleagues turns 5 interviews into 12 without additional cold outreach. Q21 converts the interview into a future demo opportunity.

Q19 Is there a specific gap in the current research tooling landscape that you wish someone would solve — something you keep running into that nothing handles well?

- This is an open invitation to hand you a product idea. Whatever they say, ask: 'How often does that come up for you?'

Q20 Who else in your world do you think would have strong opinions on this? Anyone you'd recommend I talk to?

- Always ask this. A warm referral converts at 3–4x cold outreach. If they give you a name, ask: 'Would you be comfortable introducing us over email?'

Q21 If I come back in a few weeks with something to show you — would you be open to taking a look and giving me honest feedback?

- A yes here is a soft demo commitment. Follow up within 48 hours with a calendar link. Do not wait — the window closes fast.

Facilitation Notes — How to Run These Well

ON SILENCE

After Q6 and Q12, pause after they answer and do not fill the silence immediately. Researchers are precise thinkers — their second answer is usually more revealing than their first. Count to 3 in your head before following up.

ON NEGATIVE REACTIONS

If they say 'AI pentesting is a gimmick' or express strong skepticism at Q12 or Q14, do not defend the idea. Say: 'That's really useful — what would change your mind?' Their answer tells you exactly what your product needs to prove.

ON STAYING NEUTRAL

Never say 'so you'd say the problem is X, right?' or 'most people I talk to say Y.' Leading questions corrupt your evidence. Ask open questions and repeat their words back to confirm, not to guide: 'You mentioned the handoff breaks down — can you say more about that?'

ON QUOTING

Your strongest Log 4 evidence is verbatim quotes. Write them down exactly — paraphrasing weakens the signal and introduces your own interpretation. If recording, timestamp the moments where they describe pain so you can find them later.

ON THE SAFETY ORACLE

Q18 is the most product-specific question in this guide. If at least 3 of 5 researchers say they need hard safeguards or a full audit trail before trusting an autonomous system — that is direct validation that the Safety Oracle is a selling point, not just an internal guardrail. Quote it in Log 4.

ON SUCCESS CRITERIA

Log 4 hypothesis is confirmed if: at least 4 of 8 interviewees spontaneously describe a specific pain point (from Q4–Q9) without being prompted, AND at least 3 of 8 express interest in seeing the product at Q21. If you can't reach those numbers, the ICP needs adjustment.

Quick Reference — All 21 Questions

SEC	#	QUESTION (SHORT FORM)	PURPOSE
1 — Background	1	Current research focus	Orient to role
	2	Typical week + tooling	Surface tool stack
	3	Red/purple team involvement	Assess perspective depth
2 — Workflow	4	How they study attack techniques today	Core workflow discovery
	5	Most-used tools	Tool dependency map
	6	Where the process breaks down	Primary pain discovery
	7	Validating defensive controls — example	Concrete pain evidence
	8	Known vs. novel technique frequency	Automation fit signal
	9	What gets lost in handoff to blue team	Output quality gap
3 — AI Attitudes	10	Experience with automated sim tools	Baseline AI familiarity

SEC	#	QUESTION (SHORT FORM)	PURPOSE
	1 1	What tools get right/wrong vs. humans	Gap = product opportunity
	1 2	Gut reaction to fully autonomous attacker	Trust threshold
	1 3	What system needs to do to be trusted	Feature validation
	1 4	Where AI helps vs. is oversold	Objection mapping
4 — Trust	1 5	What they need in the output to believe it	Finding data model fit
	1 6	MITRE ATT&CK — signal or theater	Technique accuracy bar
	1 7	Worst outcome of a false positive	FP tolerance calibration
	1 8	Reaction to unauthorized autonomous action	Safety Oracle validation
5 — Closing	1 9	Gap they wish someone would solve	Whitespace discovery
	2 0	Who else to talk to	Referral generation
	2 1	Open to seeing the product later	Demo pipeline